

Research and Build a SysSecOps System Process for a Hybrid Cloud Environment

Nguyễn Chính Thông — 22125102

Huỳnh Tuấn Minh — 22125055

Advisors: PhD. Tran Trung Dung · MSc. Chung Thuy Linh

Advanced Program in Computer Science

University of Science, VNU-HCM — 2026

Roadmap

1. **Problem** — why hybrid cloud operations break security
2. **Proposed process** — a SysSecOps workflow with a gate before deploy
3. **Risk engine** — turning many tools into one explainable decision
4. **Prototype** — what was actually built
5. **Results** — three research questions, three bodies of evidence
6. **Limitations & future work**

Motivation: hybrid cloud raises the stakes

- Organisations combine **public-cloud elasticity** with **private-infrastructure control**
- Hybrid is now the default operating model, not an edge case
- But the combination **multiplies operational surface**:
 - two provisioning mechanisms, two security models, two identity domains
 - configuration drift and misconfiguration become the dominant failure mode
- Provisioning, administration, delivery and security assessment are typically run by **separate teams with separate tools**

| The flexibility is real — so is the coordination cost.

Four operational challenges

	Challenge	Consequence
C1	Managing heterogeneous infrastructure	Inconsistent policy across on-prem, private and public cloud
C2	Integrating continuous security into the lifecycle	Security applied late → expensive, delayed remediation
C3	Assessing findings and prioritising risk	Many tools, many formats, no consistent prioritisation
C4	Establishing an integrated SysSecOps process	SysAdmin / Dev / Sec silos, communication gaps, no traceability

Research gaps in the literature

- **G1 — Limited operational frameworks for hybrid cloud security**
NIST RA defines *what* the components are, not *how* to operate them together
- **G2 — Continuous security stops at CI/CD**
Rarely extended across infrastructure operations in heterogeneous environments
- **G3 — Little orchestration of heterogeneous security tools**
Research categorises tools; far less work consolidates their outputs into one decision
- **G4 — Limited empirical validation of integrated processes**
Conceptual architectures dominate; executable, evaluated processes are scarce

Myrbakken & Colomo-Palacios (2017) · Rajapakse et al. (2021) · Zhao et al. (2024) · Liu et al. (2011)

Contributions

C1 — A SysSecOps operational process for hybrid cloud

Integrates system administration, continuous security and software delivery into one workflow, with a *common decision path* across public and private targets.

C2 — A unified security assessment mechanism

Normalises and deduplicates findings from multiple heterogeneous tools, combines them with cost, live compliance and a learned code-risk signal into one explainable score.

C3 — A working prototype and reproducible evaluation framework


End-to-end implementation plus separate acceptance criteria per research question, so a successful command or an intentional rejection is never mistaken for evidence of a broader claim.

Core idea: gate before deploy


- The decision is taken on the **synthesized artefact** — the CloudFormation template or the resolved Ansible play — not on the prompt or the high-level source
→ the last representation before any resource exists, and the first point where every resource and property is fully known
- **One engine for the whole hybrid:** identical weights, deduplication rules and thresholds on both sides of the trust boundary

Design goals: gate-before-deploy · one engine · graceful degradation · auditability · feedback-driven remediation · observability

Three-zone reference architecture



End-to-end artefact flow



Three stable data contracts hold the system together

- **Gate report** — `score`, `decision`, `components`, deduplicated `findings`, `scanner_status`
- **Execution result** — every external command returns `{return_code, output}`
- **Gate-decision event** → SSM Parameter Store + EventBridge

Hybrid deployment topology

- **Public side** — AWS: CDK → CloudFormation (VPC, EC2, S3, RDS), SSM, EventBridge, CloudWatch, CloudTrail, Lambda ops-loop, SNS
- **Private side** — on-premises Linux nodes configured by **Ansible**, registered via **SSM hybrid activation**
- **No public inbound administrative access**: the control plane traverses the encrypted link

Connectivity is a pluggable concern — mesh overlay VPN (Tailscale/WireGuard) · AWS Site-to-Site VPN (IPsec) · Direct Connect · Transit Gateway · SSM over public endpoints

The choice of link changes nothing in the gate, the scoring or the governance behaviour.

Multi-scanner assessment and normalisation

Tool	Signal	Dimension
<code>cfn-lint</code>	CloudFormation validity & best practice	misconfiguration
Checkov	Policy-as-code findings on templates	misconfiguration
<code>ansible-lint</code>	Best-practice & security findings on plays	misconfiguration
Regex secret scanner	Hard-coded credentials	misconfiguration
Infracost	Estimated monthly cost / delta	cost
AWS Config	Count of non-compliant live rules	compliance
Bandit + Semgrep → logistic regression	P(code is insecure)	code risk

Deduplication key: `(resource_id, template, category)` — colliding findings merge, keeping the highest severity and unioning the source labels.

The additive risk score

$$S = \underbrace{\sum_{f \in F} w(\text{sev}(f))}_{\text{severity}} + \underbrace{c(\Delta)}_{\text{cost}} + \underbrace{5v}_{\text{compliance}} + \underbrace{\text{round}(p \cdot 20)}_{\text{ML code risk}}$$

Severity	Points		Decision	Band	Action
Critical	20		Pass	$S \leq 20$	Auto-deploy
High	10		Review	$21 \leq S \leq 80$	Manual approval required
Medium	5		Reject	$S > 80$	Do not deploy; regenerate
Low	1				

Cost is **banded**, not continuous — this keeps the score stable against estimation noise. The additive form is what makes the report **explainable**: every point is attributable to a component.

Machine-learning code-risk component

- **Corpus** — positives from **SecurityEval** (CWE-labelled insecure Python); negatives sampled from maintained projects (`click` , `rich` , `typer`)
- **Features** — 11-dimensional count vector: Bandit severity ×3, Bandit confidence ×3, Semgrep severity ×3, two totals
- **Model** — logistic regression, 80/20 stratified split, standardised features, balanced class weights
- **Inference/training parity** — same analysers, same grouping, **same persisted scaler**
- $p = \max_{k \in \text{files}} \sigma(\mathbf{w}^\top \mathbf{x}_k + b)$ — worst file wins, so one risky file cannot be masked by many benign ones
- Unavailable model or analyser → `skipped` , contributes **0 points**

Scope statement: the model estimates **detector-visible** insecurity, not semantic flaws invisible to Bandit and Semgrep.

Prototype

Module	Responsibility
generation/	Provider-abstracted code generation (Bedrock / OpenRouter), CDK regeneration loop
security_gate/	Scanner adapters, deduplication, risk scoring, report assembly
execution/	Subprocess backbone with a stable result contract
pipeline/	synth-gate-diff-deploy orchestration, approval audit trail, credentials
monitoring/	CloudTrail, CloudWatch, ops-loop Lambda, on-prem SSM registration
risk_scoring/	Training and evaluation of the code-risk model
ui/ · scripts/	Streamlit operator console · CLI entry points

The console and the CLI invoke **the same orchestration logic** — a GUI operation follows the identical generate → assess → approve → deploy sequence.

Evaluation methodology — a pre-registered campaign

RQ	Method	Measures
RQ1 Decision conformance and enforcement correctness	30 declared scenarios, 106 offline runs , incl. 10 that remove one scanner	Decision conformance, enforcement invariants, score deltas, stage latency
RQ2 Cross-boundary enforcement equivalence	Same engine on an on-prem Linux node via mesh VPN; 11 catalogued weakness classes expressed twice, measured before and after remediation	Checkpoint attainment, detection coverage across the boundary
RQ3 Unified risk-evaluation capability	Held-out classifier evaluation + controlled boundary cases	Accuracy, P/R/F1, ROC-AUC, score & decision correctness

- **Declared first.** Scenarios, expected decisions and replicate counts were fixed in a version-controlled specification *before* any reported run executed; every statistic is computed from the persisted records by one analysis program
- **30 calibrated fixtures** — 6 band + **12 matched pairs** (11 weakness classes + 1 compliant *specificity control*). Live pricing and live account state are disabled: in the pilot they moved a *reject* fixture from 114 to 119
- **A security rejection is not a pipeline failure** — a valid *reject* that blocks deployment is *correct* execution. Denominators are reported in place of confidence intervals

RQ1 — Decision conformance and enforcement correctness

Enforcement invariant	n	Held	Rate
Run produces a persisted summary	106	106	1.00
Gate reports a status for every scanner	106	106	1.00
Unapproved <i>review</i> halts	43	43	1.00
Approved <i>review</i> writes an approval record	10	10	1.00
<i>Reject</i> blocks deployment	16	16	1.00
<i>Reject</i> writes a rejection record	16	16	1.00

Component removed	Full	Degraded	Δ	Decision change
Checkov (from <i>review</i> , CDK)	54	9	45	review → pass
Checkov (from <i>reject</i> , CDK)	114	64	50	reject → review
Secret scan (from <i>reject</i> , Ansible)	155	75	80	reject → review
ML risk model (from <i>reject</i> , CDK)	114	110	4	reject → reject

RQ2 — Cross-boundary equivalence: checkpoint attainment

Checkpoint	Offline (59)	Live target (6)
Target reachability	—	6/6
Syntax validation	59/59	6/6
Security gate	59/59	6/6
Dry run	31/31	3/3
Apply	—	3/3
Post-apply verification	—	3/3
Repeated-run idempotency	—	3/3

Finding 2 — one decision function, unequal detection

11 classes derived from MITRE CWE-1008, not chosen by the author: 223 members → three mechanical filters → 38 candidates → the 11 expressible as a matched pair. Measured **twice** — at a frozen gate commit, then after remediation.

Arm	Cloud	On-prem	Named by both
Baseline — gate frozen at 6e092b9	7/11	5/11	3/11
Remediated — +9 rule families	10/11	11/11	10/11

Baseline discordance: cloud-only **4** · on-prem-only **2** · neither **2** → a gate treating the branches as interchangeable would be wrong about **8 of 11**

RQ3 — Unified risk-evaluation capability

Held-out classifier performance ($n = 31$: 16 negative, 15 positive)

Accuracy	Precision	Recall	F1	ROC-AUC
0.8710	0.8667	0.8667	0.8667	0.8792

Confusion matrix	Predicted –	Predicted +		Top coefficients	
Actual –	14	2		total_semgrep	1.136
Actual +	2	13		bandit_conf_medium	1.004
				semgrep_high	0.915

Controlled gate-decision check: $S = 20 \rightarrow \text{pass}$ · $S = 21 \rightarrow \text{review}$ · $S = 80 \rightarrow \text{review}$ · $S = 81 \rightarrow \text{reject}$ · duplicate finding across tools → **counted once**. All five cases agreed with the specification.

Limitations and threats to validity

1. **Self-authored fixtures, limited scale** — pre-registration removes selection freedom, not the dependency: the gate's own authors wrote the fixtures. Ten approved-review runs, sixteen rejections, three replicates per deployment checkpoint
2. **Narrow deployment-side evidence** — apply, verification and idempotency measured against *one* fresh VM, one OS image: nothing on host heterogeneity, concurrency or partial failure
3. **Restricted implementation scope** — Python + Bandit/Semgrep; AWS CDK/CloudFormation + Ansible. No claim of transfer to other languages, providers or vulnerability classes
4. **Dataset and model validity** — 31 held-out samples from one split; undetectable categories excluded; negatives are *reference secure*, not verified secure
5. **Risk assumptions and the fail-open score** — weights, bands and thresholds uncalibrated; missing evidence is read as low risk
6. **Unequal detection coverage** — a shared decision function over unshared scanner sets. The class list is external, but the fixtures are hand-authored, so the figures bound **rule-to-CWE alignment**, not field recall
7. **Limited comparative measures** — no manual baseline; no operator effort, lead time, change-failure rate or remediation time

Conclusion

- Hybrid cloud fragments provisioning, administration, delivery and security; this work makes **security an explicit, auditable operational decision** inside that lifecycle
- **One process, one engine, both sides** of the hybrid — the same weights, deduplication and thresholds govern a CloudFormation template and an Ansible play
- The **additive score** converts heterogeneous tool output into a single decision *without* losing the per-component explanation
- The pre-registered campaign shows the implemented gate **decides and enforces as specified**: 106/106 conformance, six enforcement invariants at rate 1.00
- It also shows the **specification itself permits a release when evidence is missing**, and that a shared decision function does not by itself produce equal treatment when detection coverage differs
- Those two findings are **as much a result of this work as the conformance figures** — internal correctness under pre-declared conditions, not operational effectiveness at scale

Future work

1. **Close the fail-open condition** (*first priority*) — make the decision function consume the scanner-status record it already writes: an assurance penalty proportional to missing coverage, or a fail-closed policy forcing *review* below a coverage threshold. The existing degradation scenarios are ready-made test cases
2. **Close the last coverage gap at source** — give the cloud path a secret scan that separates an *embedded* credential from the idioms that legitimately *handle* one (entropy + provider formats, not keywords); give the private path an equivalent learned signal or an explicit *not applicable*; replace the hand-authored fixtures with independently written ones
3. **Broaden deployment-side evidence** — divergent prior state, concurrent targets, induced partial failures, so failure attribution has something other than `none` to classify; independently authored fixtures
4. **Validate the model** — larger multi-project corpus, project-aware partitioning, repeated cross-validation, external holdout, and a retained *challenge set* of analyser-invisible flaws
5. **Calibrate the score** — expert judgement and historical incidents; sensitivity and ablation against severity-only and scanner-only baselines
6. **Extend platform coverage** — more providers, languages and orchestrators; an operator study on whether the breakdown genuinely aids remediation

Thank you

Questions and discussion

Nguyễn Chính Thông · Huỳnh Tuấn Minh

Advisors: PhD. Tran Trung Dung · MSc. Chung Thuy Linh

Backup — Connectivity options

Method	Mechanism	Trade-offs
Mesh overlay VPN (Tailscale/WireGuard)	Userspace WireGuard, NAT traversal, identity ACLs	Fast, no inbound ports; depends on an agent + coordination service
AWS Site-to-Site VPN (IPsec)	Managed tunnels to a virtual private / transit gateway	Standards-based; needs a gateway device, bandwidth-limited per tunnel
AWS Direct Connect	Dedicated private circuit	Low latency, stable bandwidth; cost and lead time, usually paired with a VPN
Transit Gateway / VPN CloudHub	Managed hub-and-spoke routing	Scales multi-site routing; adds a managed tier and cost
SSM hybrid activation over public endpoints	Agent reaches SSM endpoints over TLS 443	Simplest management-plane reach; no private data-plane connectivity

Backup — Full model coefficients

Feature	Coefficient	Feature	Coefficient
total_semgrep	1.136296	bandit_high	0.435620
bandit_conf_medium	1.003949	bandit_low	0.261740
semgrep_high	0.914679	bandit_conf_high	0.138912
semgrep_medium	0.720708	bandit_conf_low	0.083852
total_bandit	0.610932	semgrep_low	0.067123
bandit_medium	0.454592		

Features were standardised before training, so magnitudes indicate relative influence **within this fitted model**. Total-finding features are mathematically related to the severity counts, introducing correlated predictors — the coefficients describe the model, not causation.

Backup — Dataset construction

- **Positive class** — SecurityEval, CWE-labelled insecure Python, one file per sample, label 1
- **Negative class** — files sampled from `click`, `rich`, `typer`, label 0 (*reference secure*)
- **Filtering** — exclude package initialisers and tests; admit files with 20–700 non-blank, non-comment lines (model corpus narrowed to 20–300 lines to match SecurityEval scale)
- **Exclusion** — vulnerability categories not representable by the configured Bandit/Semgrep features were removed, keeping $\approx 15\%$ as hard negatives
- **Effect** — an initial model over the unfiltered corpus reached only ≈ 0.73 accuracy with recall ≈ 0.46 ; the filtered corpus makes the chosen features informative
- **Result** — ≈ 150 samples, near-balanced; 31-file held-out partition

Backup — Comparison with related approaches

Approach	Primary emphasis	Relationship to this work
Traditional operations	Separately coordinated provisioning, administration, review	This work integrates the transitions; no quantitative manual baseline yet
NIST Reference Architecture	Vendor-neutral actors and components	Supplies context; this work adds a concrete but less vendor-neutral sequence
CloudCAMP	Model-driven generation, less manual scripting	Stronger platform abstraction; this work emphasises gating and traceability
Industrial IaC practice	Adoption, integration, maintenance challenges	Normalisation addresses heterogeneous output but creates an adapter layer to maintain
Continuous security (ADOC, Rajapakse et al.)	Lifecycle security, tool integration, triage	This work adds hybrid decision control and an explainable quantitative score